

System Programming

Prof. Chuan-Ju Wang
Dept. of Computer Science
University of Taipei

File Systems

- A disk can be divided into **logical partitions**.
- Each logical partition may be further divided into cylinder groups/partitions containing filesystems.
- Each filesystem contains a list of inodes (i-list) as well as the actual directory and data blocks.

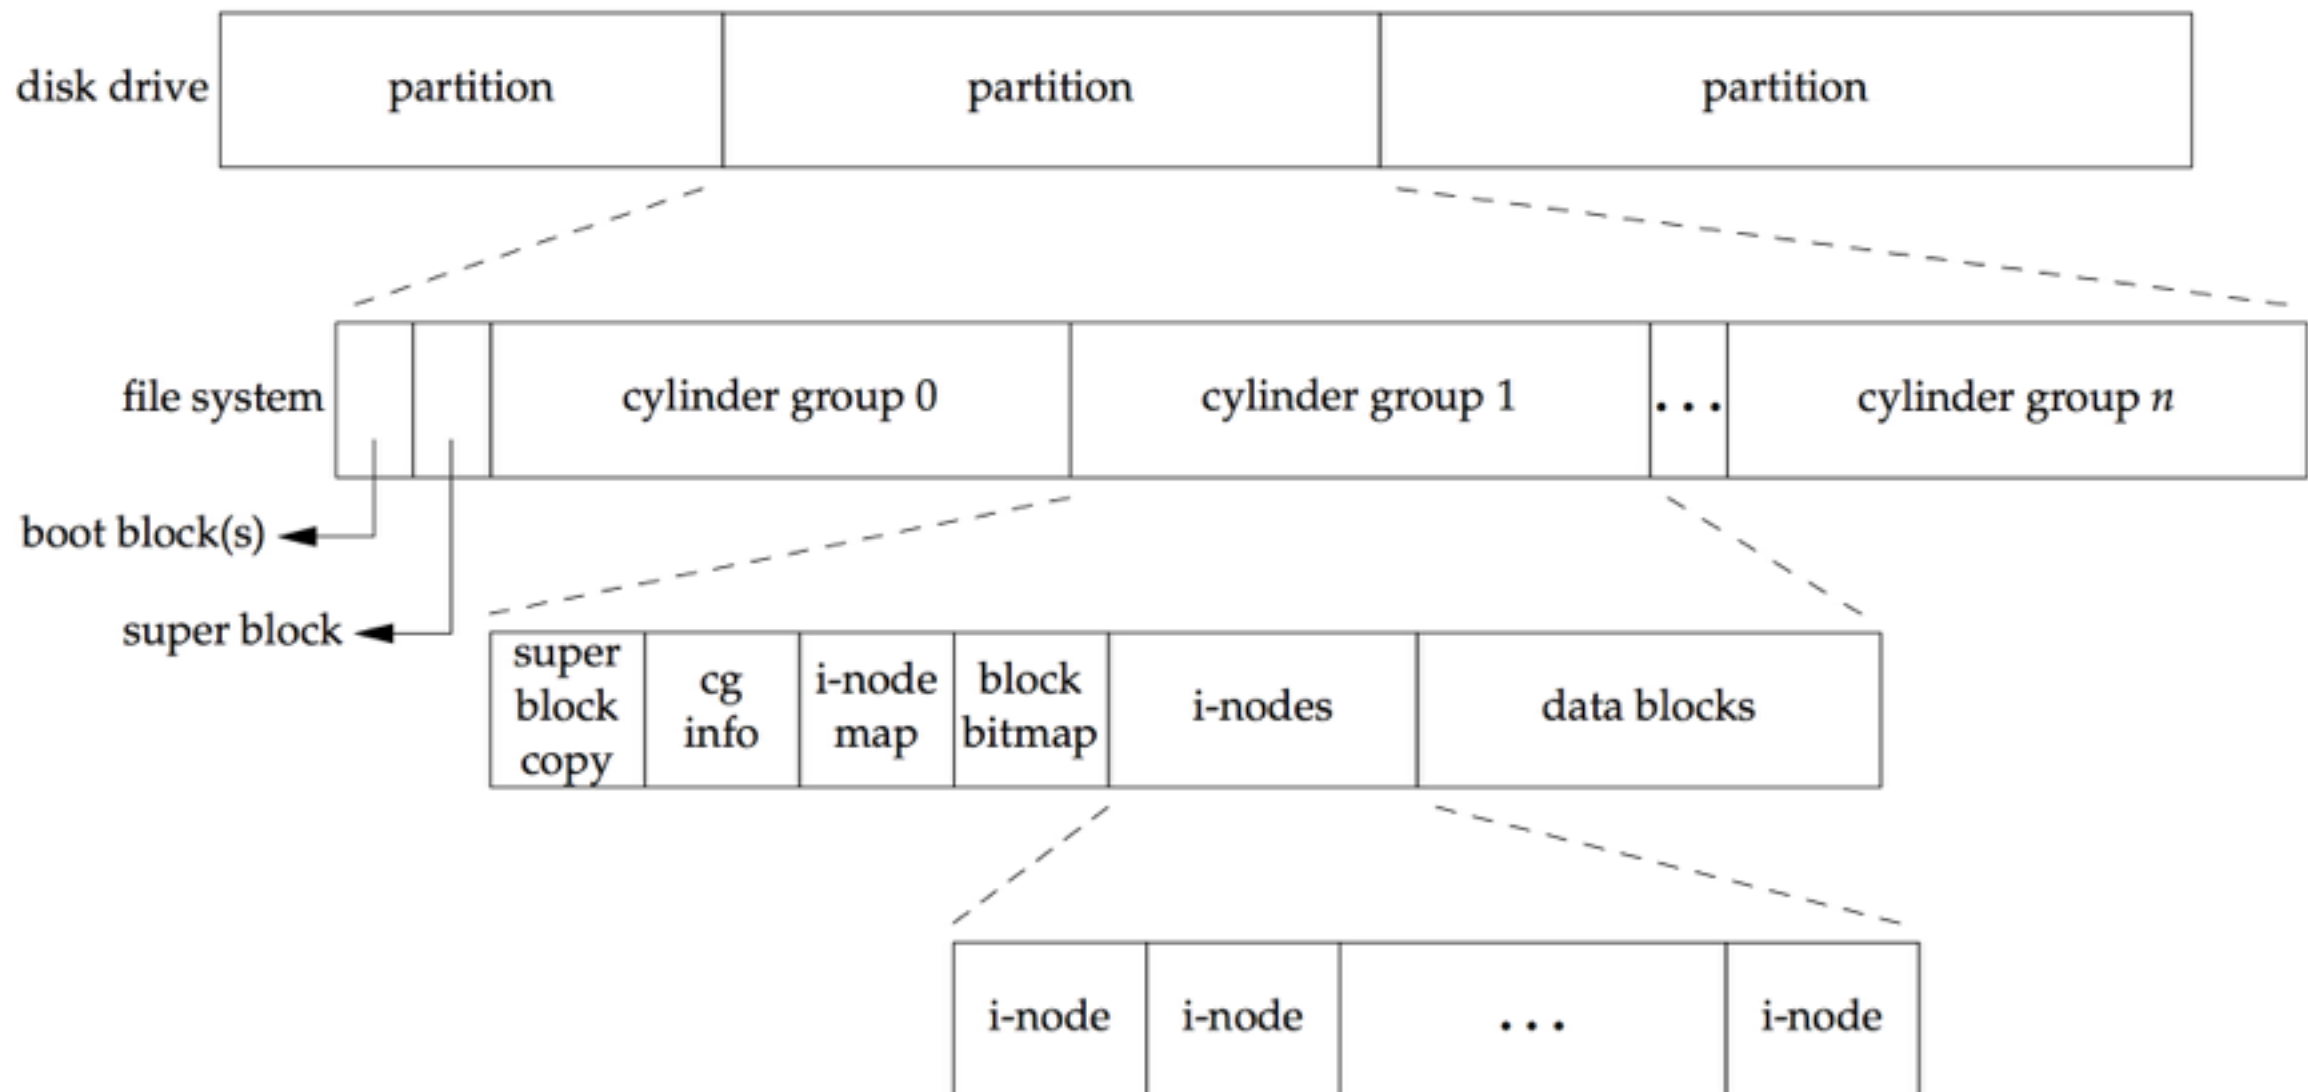


Figure 4.13 Disk drive, partitions, and a file system

- i-nodes*: **fixed-length entries** that contain most of the information about a file.
 - Eg., The file type, the file's access permission bits, the size of the file, pointers to the file's data blocks...

* <http://en.wikipedia.org/wiki/Inode>

- A directory entry is really just a **hard link** mapping a “filename” to an i-node.
- You can have many such mappings to the same file.

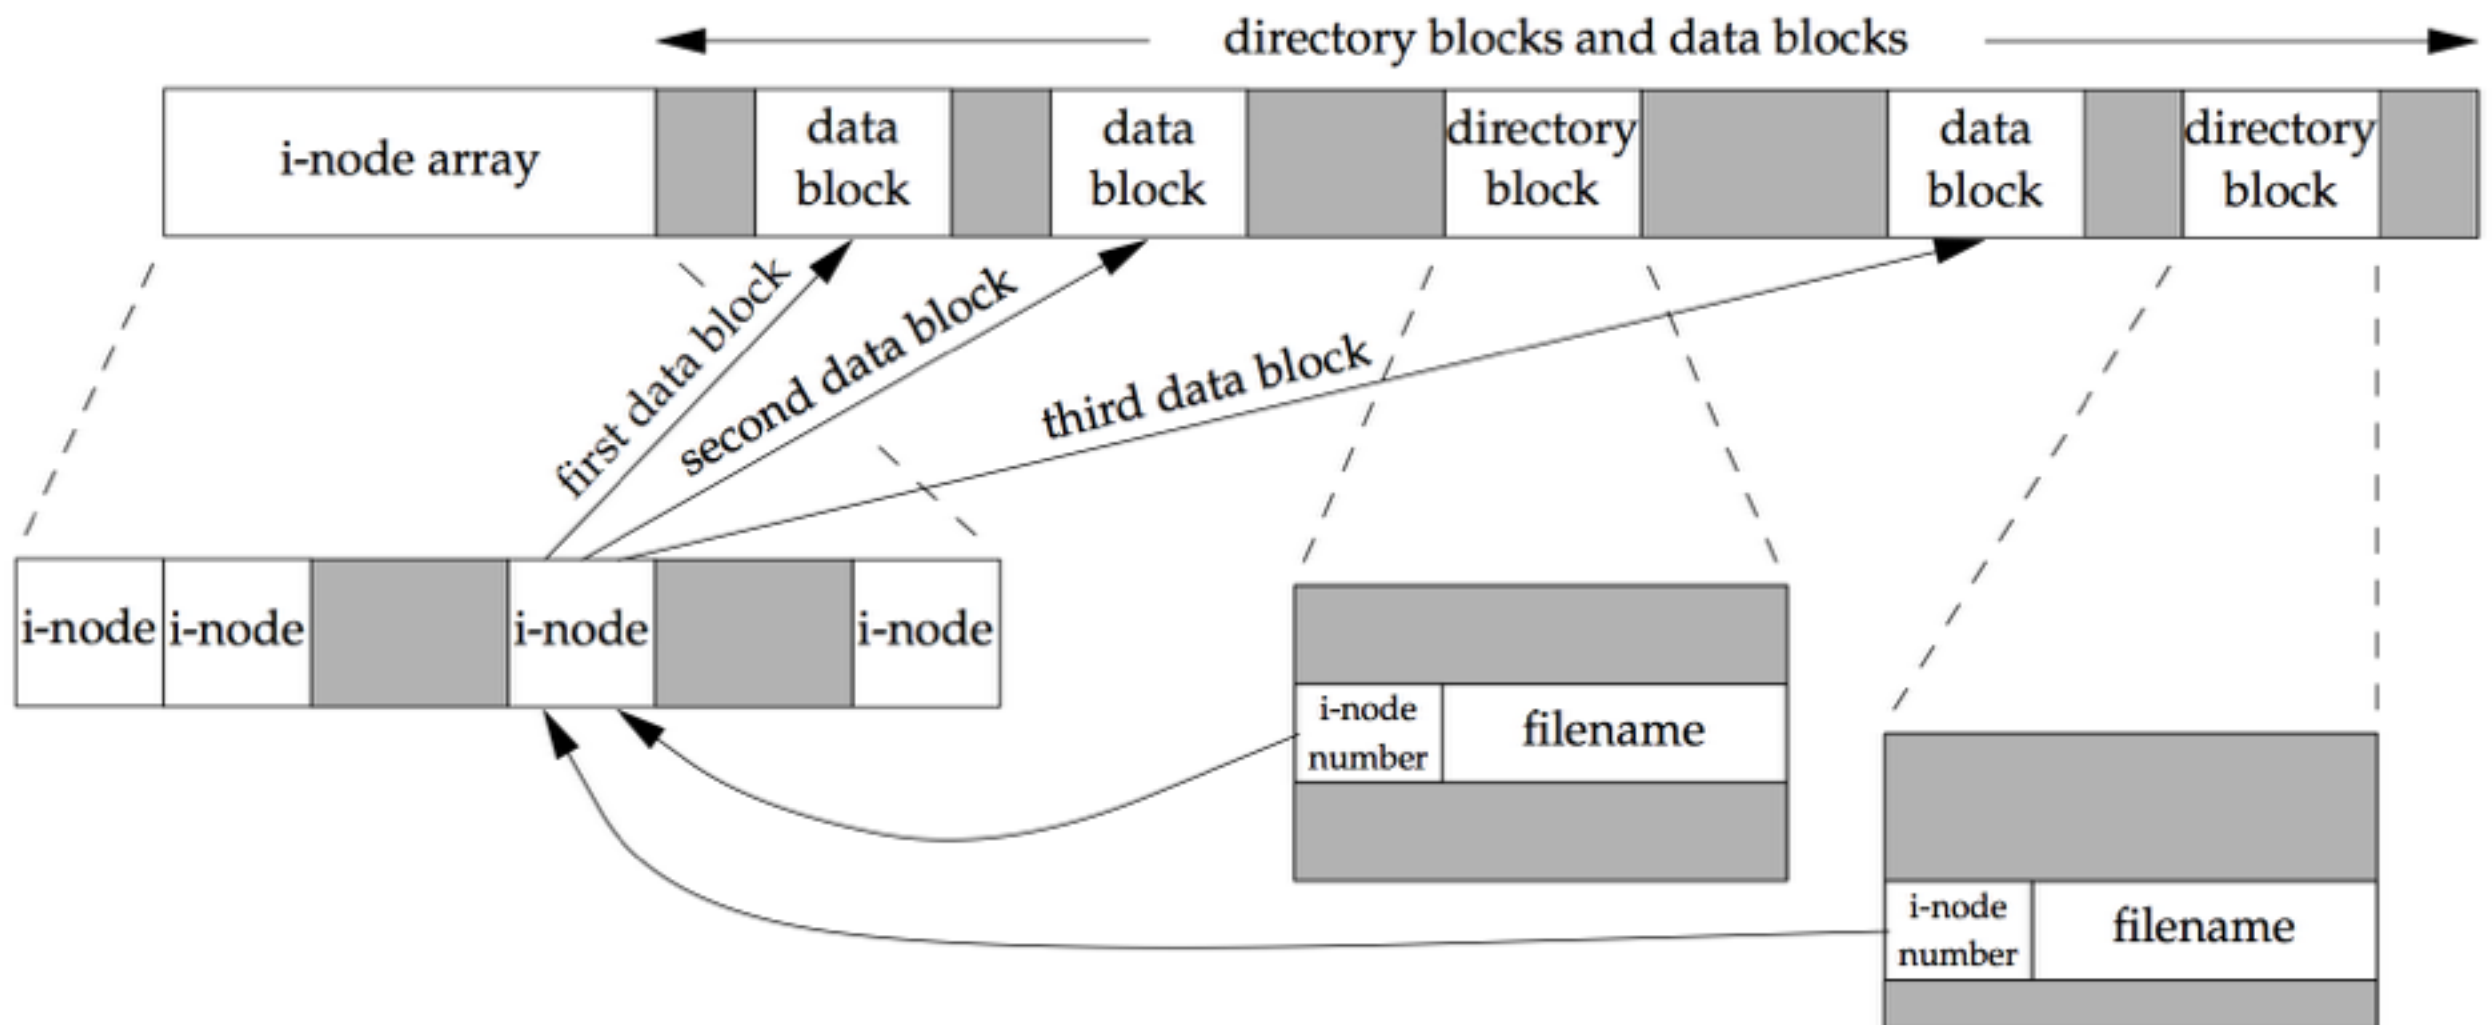


Figure 4.14 Cylinder group's i-nodes and data blocks in more detail

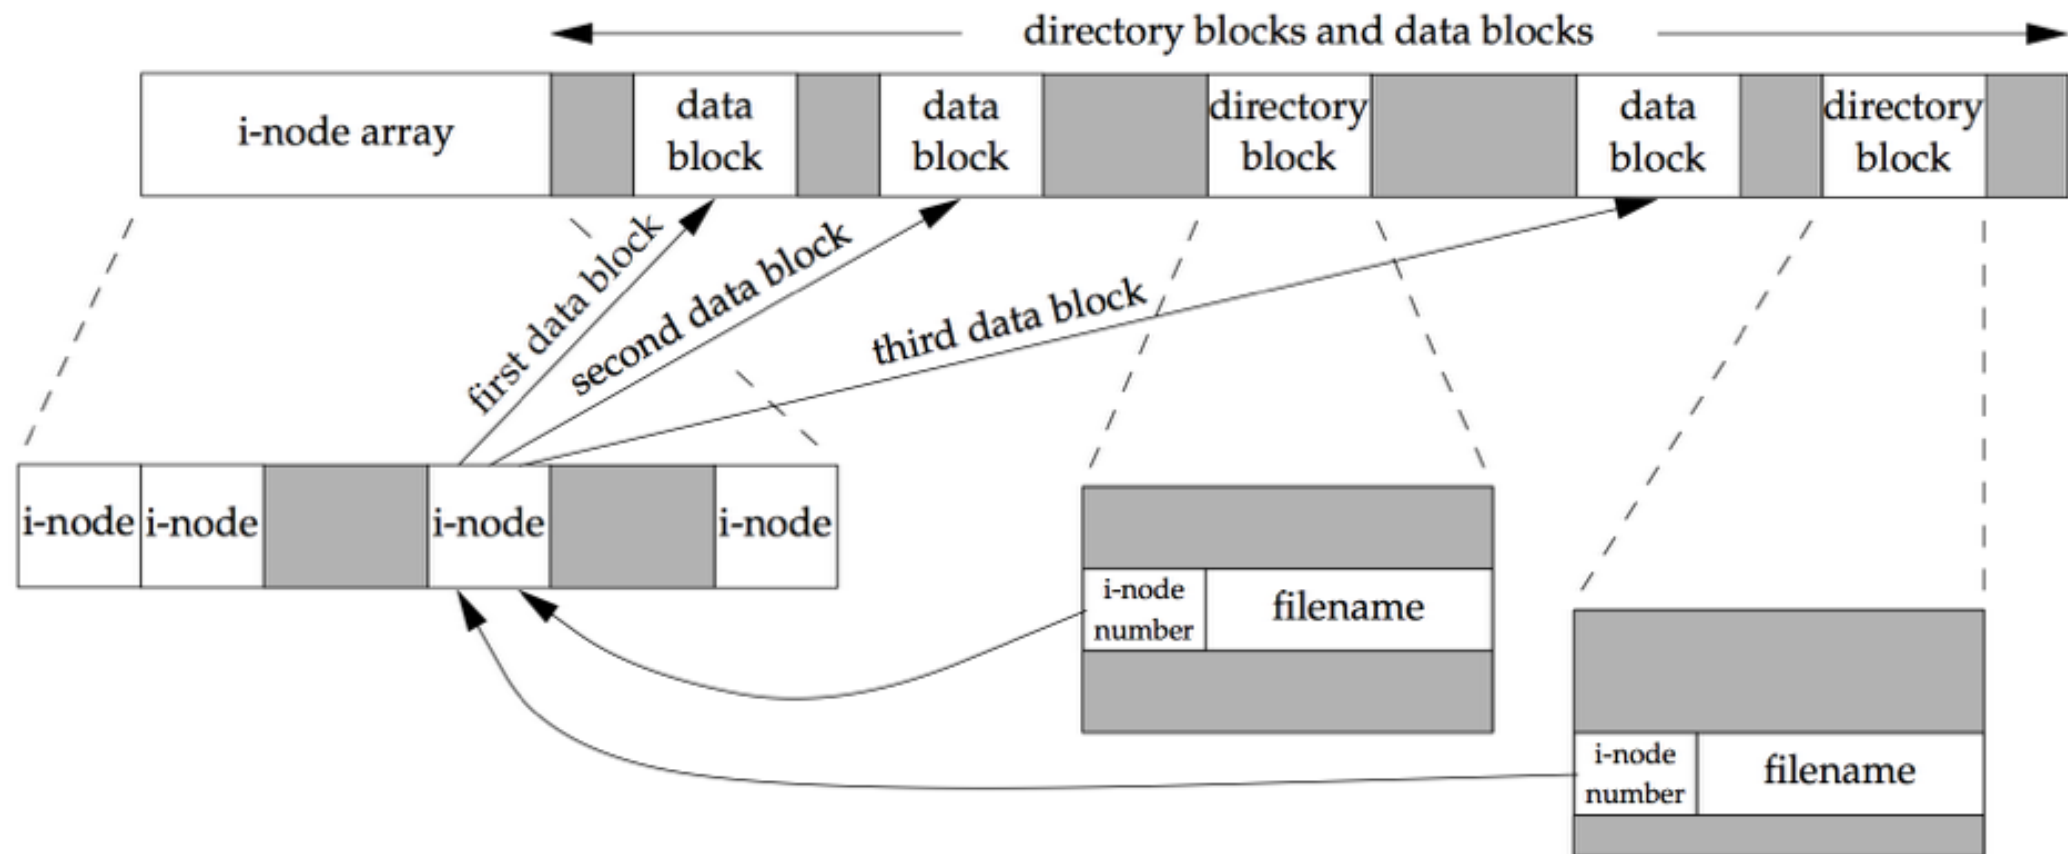


Figure 4.14 Cylinder group's i-nodes and data blocks in more detail

- Every i-node has a **link count** that contains the number of directory entries that point to the i-node.
- Only when the link count goes to 0 can the file be deleted.
 - UNLINK a file does not mean always DELETE a file!!
- Link count is contained in `st_nlink` (in `stat`).

```
struct stat {
    mode_t    st_mode;        /* file type & mode (permissions) */
    ino_t     st_ino;         /* i-node number (serial number) */
    dev_t     st_dev;         /* device number (file system) */
    dev_t     st_rdev;        /* device number for special files */
    nlink_t    st_nlink;      /* number of links */
}
```

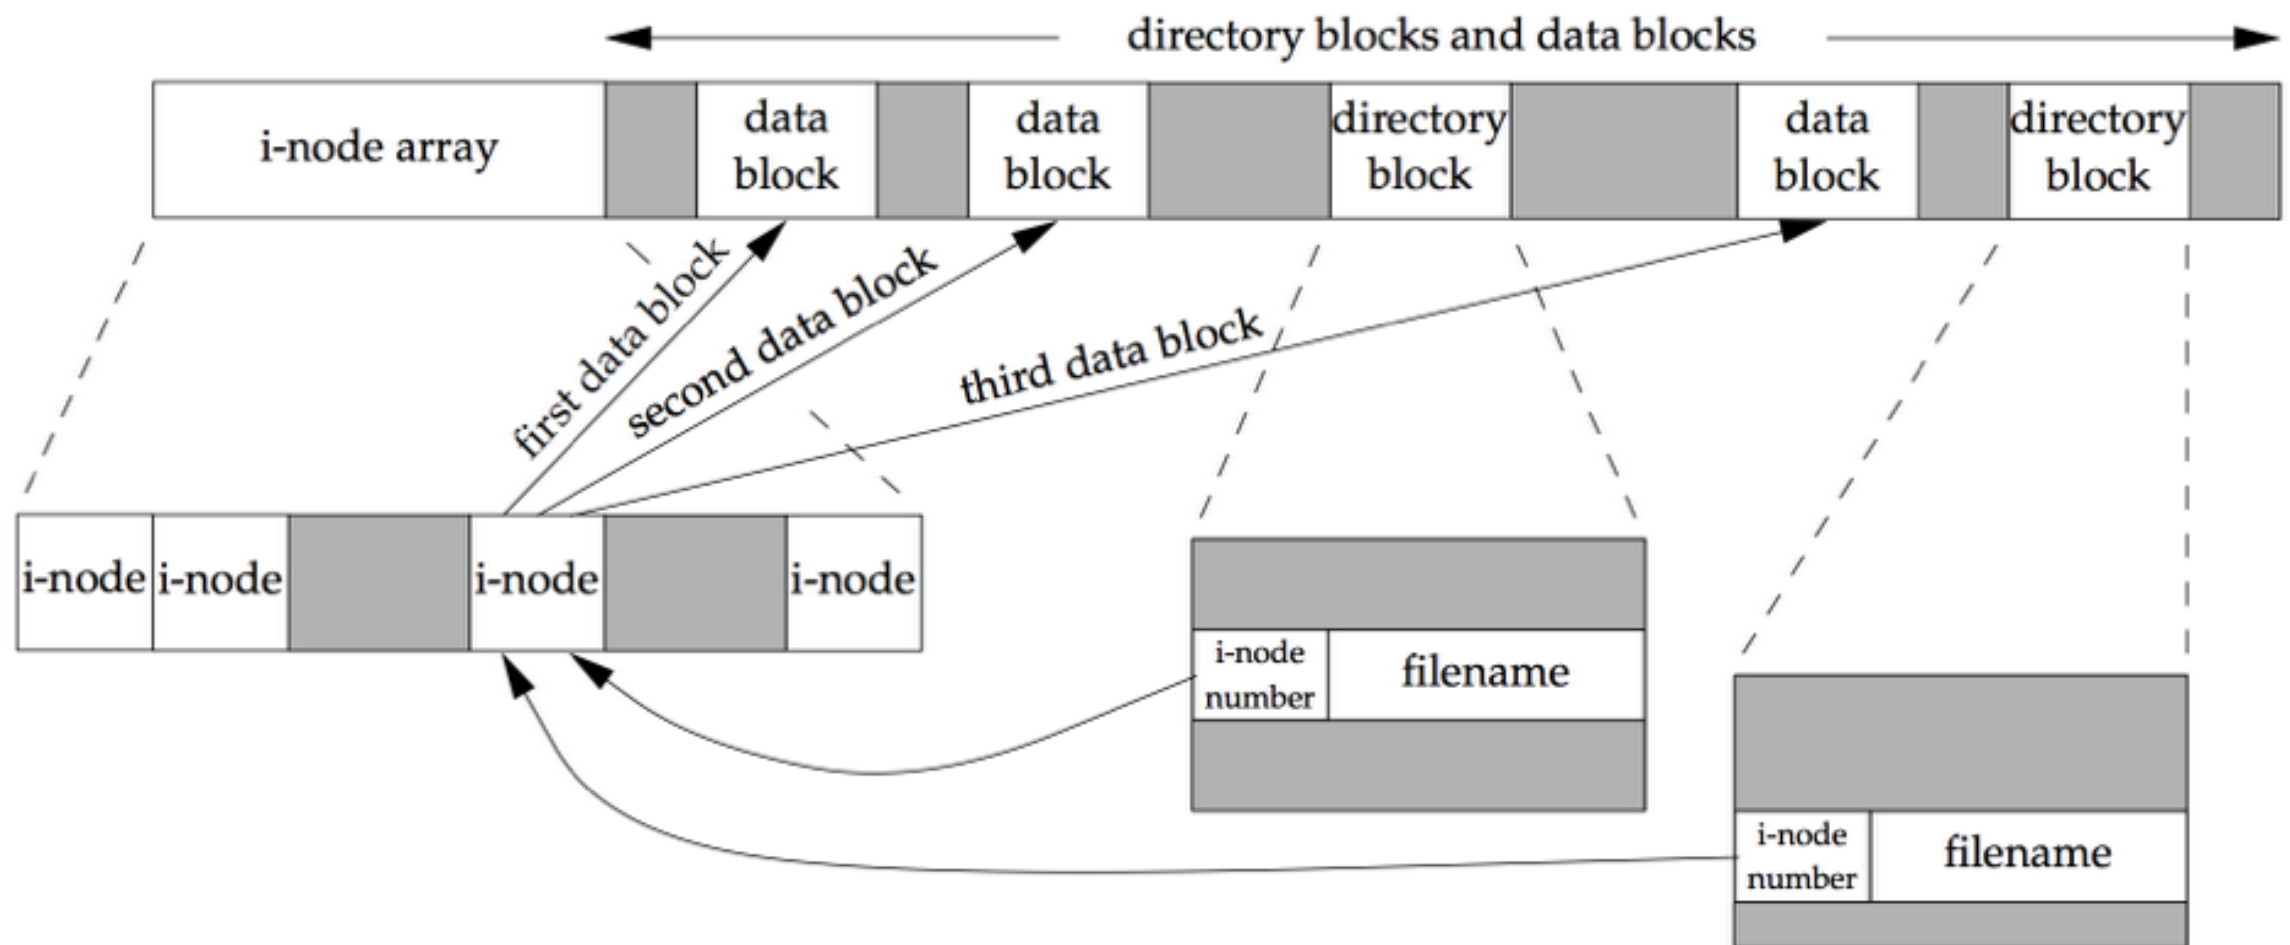


Figure 4.14 Cylinder group's i-nodes and data blocks in more detail

- **i-node number** in a directory entry must point to an i-node on the same file system (**no hard links across filesystems**).
- To move a file within a single filesystem, we can just "move" the directory entry (actually done by **creating a new entry, and deleting the old one**).

File Systems - Directories

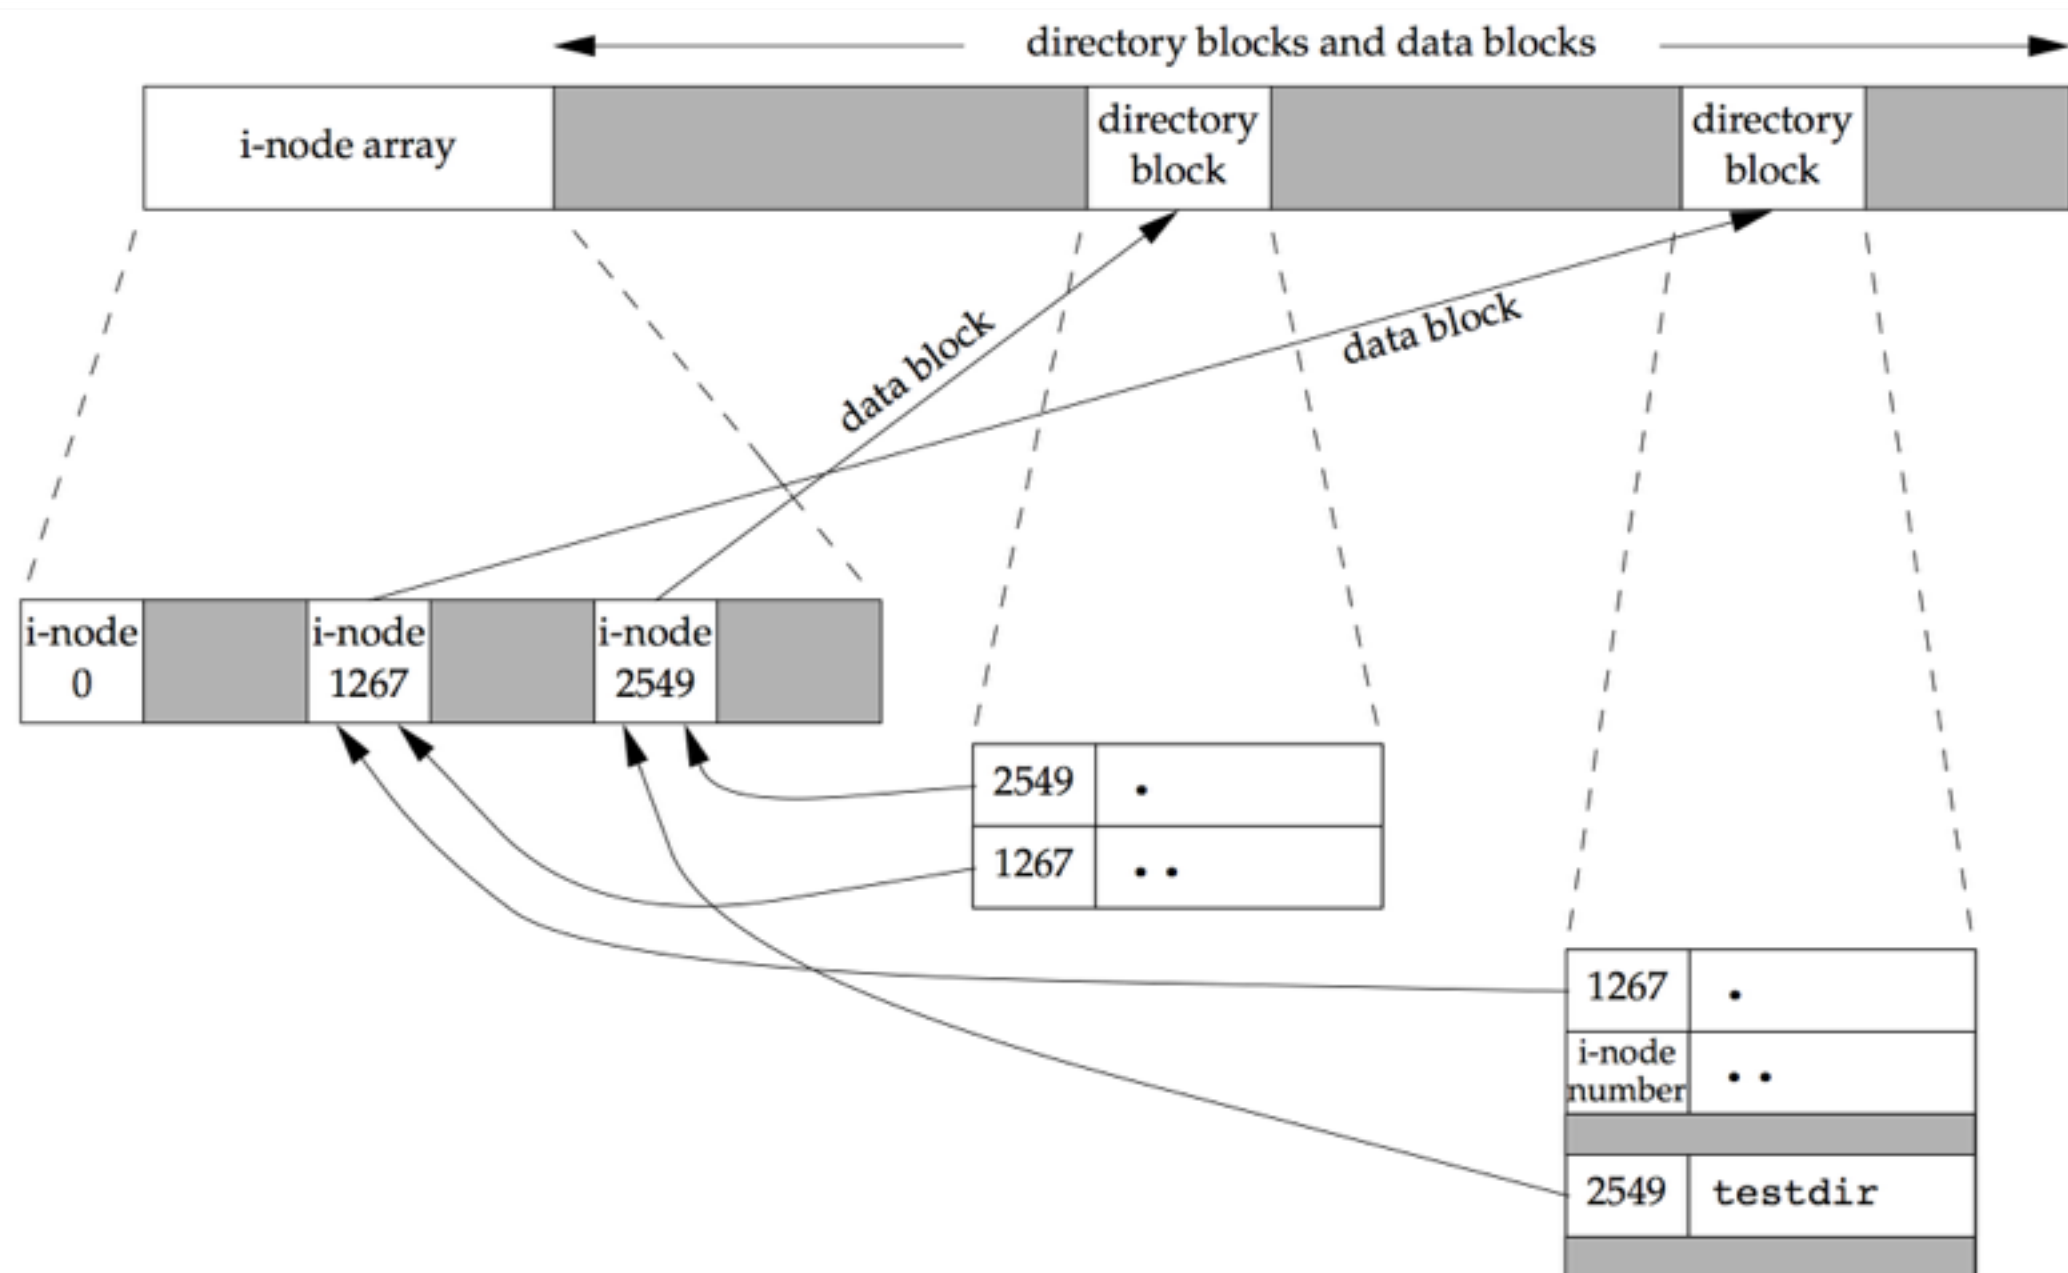


Figure 4.15 Sample cylinder group after creating the directory `testdir`

File Systems - Directories

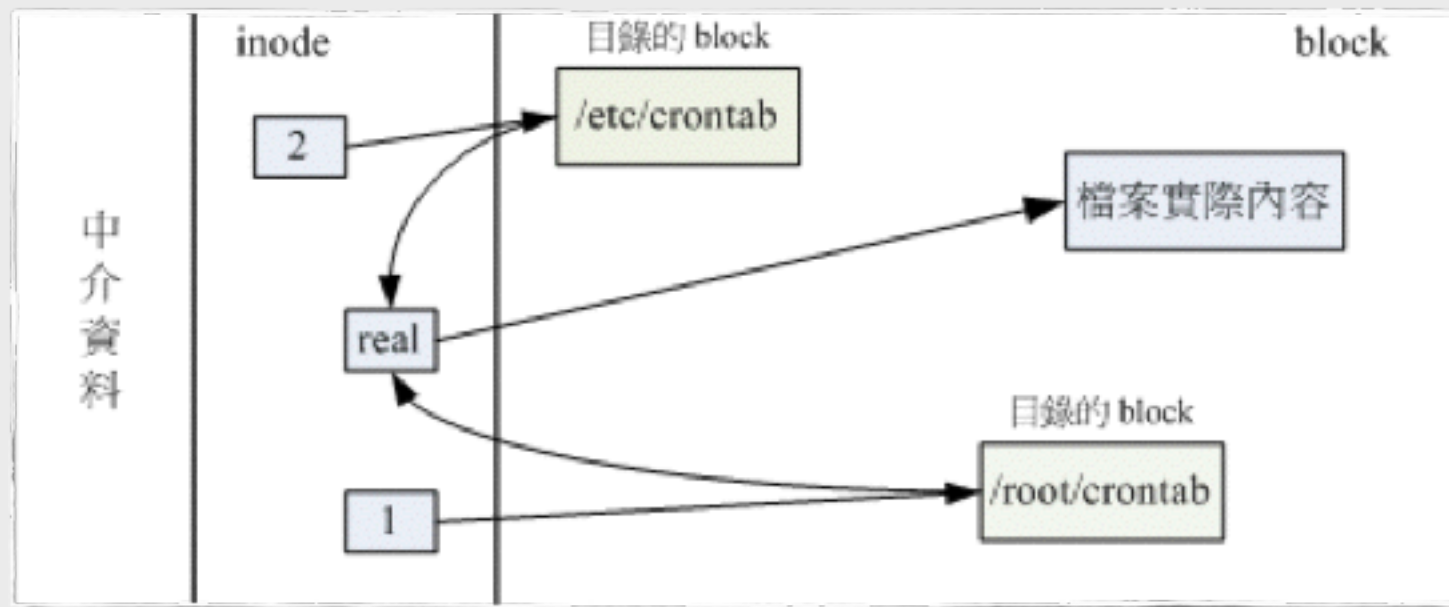
- Directories are special "files" containing hard links.
- Each directory contains at least two entries:
 - `.` (this directory)
 - `..` (the parent directory)
- The link count (`st_nlink`) of a directory is **at least 2***

* http://teaching.idallen.com/dat2330/04f/notes/links_and_inodes.html

Hard Links

- `ln(1)`

```
[root@www ~]# ln /etc/crontab .    <==建立實體連結的指令
[root@www ~]# ll -i /etc/crontab /root/crontab
1912701 -rw-r--r-- 2 root root 255 Jan  6  2007 /etc/crontab
1912701 -rw-r--r-- 2 root root 255 Jan  6  2007 /root/crontab
```



- Limitations”
 - Most modern operating systems don't allow hard links on **directories**.
 - Hard links can only be created to files **on the same file system**.

link(2) and unlink(2)

- link(2) function creates a new dir entry that references an existing file (hard link).

```
#include <unistd.h>

int link(const char *existingpath, const char
    ↪ *newpath);
```

Returns: 0 if OK, 1 on error

- This function creates a new directory entry, *newpath*, that references the existing file *existingpath*.
 - If the *newpath* already exists, an error is returned.
- POSIX.1 allows links to cross filesystems, most implementations (SVR4, BSD) don't.
- Only **superuser** can create links to directories (loops in filesystem are bad).

link(2) and unlink(2)

- unlink(2) function removes an existing dir entry and decrements the link count of the file.

```
#include <unistd.h>
```

```
int unlink(const char *pathname);
```

Returns: 0 if OK, 1 on error

- If there are other links to the file, the data in the file **is still accessible through the other links**.
- To unlink a file, we must have **write permission** and **execute permission** in the directory containing the directory entry

Figure 4.16. Open a file and then **unlink** it

```
#include "apue.h"
#include <fcntl.h>

int
main(void)
{
    if (open("tempfile", O_RDWR) < 0)
        err_sys("open error");
    if (unlink("tempfile") < 0)
        err_sys("unlink error");
    printf("file unlinked\n");
    sleep(15);
    printf("done\n");
    exit(0);
}
```

```
$ ls -l tempfile                                look at how big the file is
-rw-r----- 1 sar      413265408 Jan 21 07:14 tempfile
$ df /home                                     check how much free space is available
Filesystem 1K-blocks    Used Available Use% Mounted on
/dev/hda4   11021440 1956332   9065108  18% /home
$ ./a.out &                                  run the program in Figure 4.16 in the background
1364                                             the shell prints its process ID
$ file unlinked                               the file is unlinked
ls -l tempfile                                see if the filename is still there
ls: tempfile: No such file or directory        the directory entry is gone
$ df /home                                    see if the space is available yet
Filesystem 1K-blocks    Used Available Use% Mounted on
/dev/hda4   11021440 1956332   9065108  18% /home
$ done                                         the program is done, all open files are closed
df /home                                       now the disk space should be available
Filesystem 1K-blocks    Used Available Use% Mounted on
/dev/hda4   11021440 1552352   9469088  15% /home
                                             now the 394.1 MB of disk space are available
```

`link(2)` and `unlink(2)`

- `unlink` is often used by a program to ensure that a temporary file it creates **won't be left around in case the program crashes**.
- The process creates a file using either `open` or `creat` and then **immediately calls `unlink`**.
- The file is not deleted, however, because it is still open.
- Only when the process **either closes the file or terminates**, which causes the kernel to close all its open files, the file is deleted.

remove (2)

- Unlink a file or a directory with the `remove` function.
- For a file, `remove` is identical to `unlink`.
- For a directory, `remove` is identical to `rmdir`.

```
#include <stdio.h>

int remove(const char *pathname);
```

Returns: 0 if OK, 1 on error

rename (2)

```
#include <stdio.h>
```

```
int rename(const char *oldname, const char *newname);
```

Returns: 0 if OK, 1 on error

- If *oldname* refers to a file:
 - If *newname* exists and it is not a directory, it's removed and *oldname* is renamed *newname*.
 - If *newname* exists and it is a directory, an error results.
 - Must have *w+x* permissions for the directories containing *old/newname*.
- If *oldname* refers to a directory:
 - If *newname* exists and is an empty directory (contains only . and ..), it is removed; *oldname* is renamed *newname*.
 - If *newname* exists and is a file, an error results.
 - If *oldname* is a prefix of *newname*, an error results.
 - Must have *w+x* perms for the directories containing *old/newname*.


Symbolic Links

- Recall that hard links are limited:
 - Cannot link to a different file system
 - Generally not allowed for a directory
 - Only the superuser can create a hard link to a directory (when supported by the underlying file system).

Symbolic Links

- `ln(1)`

```
[root@www ~]# ln -s /etc/crontab crontab2
[root@www ~]# ll -i /etc/crontab /root/crontab2
1912701 -rw-r--r-- 2 root root 255 Jan  6 2007 /etc/crontab
654687 lrwxrwxrwx 1 root root 12 Oct 22 13:58 /root/crontab2 -> /etc/crontab
```



Symbolic Links

- When using functions that refer to a file by name, we always need to know **whether the function follows a symbolic link**.
- If the function follows a symbolic link, a pathname argument to the function **refers to the file pointed to by the symbolic link**.
- Otherwise, a pathname argument **refers to the link itself**, not the file pointed to by the link.


Figure 4.17. Treatment of symbolic links by various functions

Function	Does not follow symbolic link	Follows symbolic link
access		•
chdir		•
chmod		•
chown	•	•
creat		•
exec		•
lchown	•	
link		•
lstat	•	
open		•
opendir		•
pathconf		•
readlink	•	
remove	•	
rename	•	
stat		•
truncate		•
unlink	•	

Symbolic Links

- It is possible to introduce **loops** into the file system by using symbolic links.
- Most functions that look up a pathname return an `errno` of `ELOOP` when this occurs.

```
$ mkdir foo           make a new directory
$ touch foo/a         create a 0-length file
$ ln -s ../foo foo/testdir create a symbolic link
$ ls -l foo
total 0
-rw-r--r-- 1 root root 0 Jan 22 00:16 a
lrwxrwxrwx 1 root root 5 Jan 22 00:16 testdir -> ../foo
```



```
foo
foo/a
foo/testdir
foo/testdir/a
foo/testdir/testdir
foo/testdir/testdir/a
foo/testdir/testdir/testdir
foo/testdir/testdir/testdir/a
```

```
0 Jan 22 00:16 a
6 Jan 22 00:16 testdir -> ../foo
```

Symbolic Links

- When we open a file, if the pathname passed to open specifies a **symbolic link**, open follows the link to the specified file.
- If the file pointed to by the symbolic link **doesn't exist**, open returns an **error**.

```
$ ln -s /no/such/file myfile      create a symbolic link
$ ls myfile
myfile                          ls says it's there
$ cat myfile                     so we try to look at it
cat: myfile: No such file or directory
$ ls -l myfile                   try -l option
lrwxrwxrwx 1 sar                 13 Jan 22 00:26 myfile -> /no/such/file
```

symlink(2) and readlink(2)

- A new directory entry, *sympath*, is created that points to *actualpath*.

```
#include <unistd.h>

int symlink(const char *actualpath, const char
    ↪ *sympath);
```

Returns: 0 if OK, 1 on error

- It is not required that *actualpath* exists when the symbolic link is created.
- Also, *actualpath* and *sympath* need not reside in the same file system.

symlink(2) and readlink(2)

- Because the open function follows a symbolic link, we need a way to **open the link itself and read the name in the link.**

```
#include <unistd.h>

ssize_t readlink(const char* restrict pathname,
    ↪ char *restrict buf,
               size_t bufsize);
```

Returns: number of bytes read if OK, 1 on error

- This function combines the actions of open, read, and close.
- If the function is successful, it returns **the number of bytes placed into buf.**
 - Note: *buf* are not null terminated.

File Times

- Three time fields are maintained for each file.

Figure 4.19. The three time values associated with each file

Field	Description	Example	ls(1) option
st_atime	last-access time of file data	read	-u
st_mtime	last-modification time of file data	write	default
st_ctime	last-change time of i-node status	chmod, chown	-c

- The difference between st_mtime and st_ctime:
 - The modification time is when the contents of the file were last modified.
 - The changed-status time is when the i-node of the file was last modified.
 - Changing the file access permissions, changing the user ID, ...

```
jere@VirtualBox-MBP [~/SystemProgramming/apue.2e] ls -l fig4.12
-rwxrwxrwx 1 jere jere 16 2005-05-28 22:18 fig4.12 -> file/changemod.c
jere@VirtualBox-MBP [~/SystemProgramming/apue.2e] ls -cl fig4.12
-rwxrwxrwx 1 jere jere 16 2012-01-09 16:25 fig4.12 -> file/changemod.c
jere@VirtualBox-MBP [~/SystemProgramming/apue.2e] ls -ul fig4.12
-rwxrwxrwx 1 jere jere 16 2012-02-19 19:58 fig4.12 -> file/changemod.c
```

utime(2)

- The access time and the modification time of a file can be changed with the `utime` function.

```
#include <utime.h>

int utime(const char *pathname, const struct
    utimbuf *times);
```

Returns: 0 if OK, 1 on error

```
struct utimbuf {
    time_t actime;    /* access time */
    time_t modtime;   /* modification time */
}
```

- If `times` is `NULL`, access time and modification time are set to **the current time**.
- If `times` is non-`NULL`, then times are set to **the values in the structure pointed to by `times`**.
- Note that `st_ctime` is set to **the current time** in both cases.

Figure 4.21. Example of `utime` function

```
#include "apue.h"
#include <fcntl.h>
#include <utime.h>

int
main(int argc, char *argv[])
{
    int          i, fd;
    struct stat   statbuf;
    struct utimbuf timebuf;

    for (i = 1; i < argc; i++) {
        if (stat(argv[i], &statbuf) < 0) { /* fetch current times */
            err_ret("%s: stat error", argv[i]);
            continue;
        }
        if ((fd = open(argv[i], O_RDWR | O_TRUNC)) < 0) { /* truncate */
            err_ret("%s: open error", argv[i]);
            continue;
        }
        close(fd);
        timebuf.actime = statbuf.st_atime;
        timebuf.modtime = statbuf.st_mtime;
        if (utime(argv[i], &timebuf) < 0) { /* reset times */
            err_ret("%s: utime error", argv[i]);
            continue;
        }
    }
    exit(0);
}
```

```

$ ls -l changemod times look at sizes and last-modification times
-rwxrwxr-x 1 sar 15019 Nov 18 18:53 changemod
-rwxrwxr-x 1 sar 16172 Nov 19 20:05 times
$ ls -lu changemod times look at last-access times
-rwxrwxr-x 1 sar 15019 Nov 18 18:53 changemod
-rwxrwxr-x 1 sar 16172 Nov 19 20:05 times
$ date print today's date
Thu Jan 22 06:55:17 EST 2004
$ ./a.out changemod times run the program in Figure 4.21
$ ls -l changemod times and check the results
-rwxrwxr-x 1 sar 0 Nov 18 18:53 changemod
-rwxrwxr-x 1 sar 0 Nov 19 20:05 times
$ ls -lu changemod times check the last-access times also
-rwxrwxr-x 1 sar 0 Nov 18 18:53 changemod
-rwxrwxr-x 1 sar 0 Nov 19 20:05 times
$ ls -lc changemod times and the changed-status times
-rwxrwxr-x 1 sar 0 Jan 22 06:55 changemod
-rwxrwxr-x 1 sar 0 Jan 22 06:55 times

```

Assignment

- Reading (do it at home):
 - Read and try:
 - http://linux.vbird.org/linux__basic/0210filepermission.php
 - http://linux.vbird.org/linux__basic/0230filesystem/0230filesystem.php
 - Manual pages for the functions covered
 - Stevens Chap. 4

Assignment

- Assignment 5

I. Coding

- Use `umask (2)` to set your umask to various values and see what happens to new files you create (similar to the program Fig. 4.9)
 - ➔ A report should be handed in (just copy your program to your report) and explain what happens?
- Write your own `ls` command
 - ➔ A program (submit via TA's server) should be handed in.
 - ➔ A brief report (upload via lms system) should be handed in.

Assignment

- Assignment 5

2. Finding and thinking

- a) Verify that turning off user-read permission for a file that you own denies you access to the file, even if group- or other permissions allow reading.
- b) What does the uppercase “S” mean?
- c) Where did the function `path_alloc` be defined and implemented? What did this function do and how did it work?

```
ptr = path_alloc(&size); /* our own function */  
if (getcwd(ptr, size) == NULL)  
    perror("getcwd failed");
```

➔ A report should be handed in (submit via lms).